

Ungraded lab: Permutation Feature Importance

Welcome, during this ungraded lab you are going to be perform Permutation Feature Importance on the [wine dataset](#) using scikit-learn. In particular you will:

1. Train a Random Forest classifier on the data.
2. Compute the feature importance score by permutating each feature.
3. Re-train the model with only the top features.
4. Check other classifiers for comparison.

Let's get started!

Inspect and pre-process the data

Begin by upgrading scikit-learn to the latest version:

```
In [1]: !pip install -U scikit-learn

Requirement already satisfied: scikit-learn in /usr/local/lib/python3.7/dist-packages (0.22.2.post1)
Collecting scikit-learn
  Downloading scikit_learn-0.24.2-cp37-cp37m-manylinux2010_x86_64.whl (22.3 MB)
    |████████████████████████████████████████| 22.3 MB 1.2 MB/s
Collecting threadpoolctl>=2.0.0
  Downloading threadpoolctl-2.2.0-py3-none-any.whl (12 kB)
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.7/dist-packages (from scikit-learn) (1.0.1)
Requirement already satisfied: scipy>=0.19.1 in /usr/local/lib/python3.7/dist-packages (from scikit-learn) (1.4.1)
Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib/python3.7/dist-packages (from scikit-learn) (1.19.5)
Installing collected packages: threadpoolctl, scikit-learn
  Attempting uninstall: scikit-learn
    Found existing installation: scikit-learn 0.22.2.post1
    Uninstalling scikit-learn-0.22.2.post1:
      Successfully uninstalled scikit-learn-0.22.2.post1
  Successfully installed scikit-learn-0.24.2 threadpoolctl-2.2.0
```

Now import the required dependencies and load the dataset:

```
In [2]: import numpy as np
from sklearn.datasets import load_wine

# as_frame param requires scikit-learn >= 0.23
data = load_wine(as_frame=True)

# Print first rows of the data
data.frame.head()
```

```
Out[2]:
```

	alcohol	malic_acid	ash	alcalinity_of_ash	magnesium	total_phenols	flavanoids	nonflavanoid_phenols	proanthocyanins	color_intensity
0	14.23	1.71	2.43	15.6	127.0	2.80	3.06	0.28	2.29	5.04
1	13.20	1.78	2.14	11.2	100.0	2.65	2.76	0.26	1.28	4.31
2	13.16	2.36	2.67	18.6	101.0	2.80	3.24	0.30	2.81	5.42
3	14.37	1.95	2.50	16.8	113.0	3.85	3.49	0.24	2.18	7.80
4	13.24	2.59	2.87	21.0	118.0	2.80	2.69	0.39	1.82	4.69

This dataset is made up of 13 numerical features and there are 3 different classes of wine.

Now perform the train/test split and normalize the data using `StandardScaler`:

```
In [3]: from sklearn.model_selection import train_test_split
        from sklearn.preprocessing import StandardScaler

        # Train / Test split
        X_train, X_test, y_train, y_test = train_test_split(data.data, data.target, random_state=42)

        # Instantiate StandardScaler
        scaler = StandardScaler()

        # Fit it to the train data
        scaler.fit(X_train)

        # Use it to transform the train and test data
        X_train = scaler.transform(X_train)

        # Notice that the scaler is trained on the train data to avoid data leakage from the test set
        X_test = scaler.transform(X_test)
```

Train the classifier

Now you will fit a **Random Forest classifier** with 10 estimators and compute the mean accuracy achieved:

```
In [4]: from sklearn.ensemble import RandomForestClassifier

        # Fit the classifier
        rf_clf = RandomForestClassifier(n_estimators=10, random_state=42).fit(X_train, y_train)

        # Print the mean accuracy achieved by the classifier on the test set
        rf_clf.score(X_test, y_test)
```

Out[4]: 0.9111111111111111

This model achieved a mean accuracy of 91%. Pretty good for a model without fine tuning.

Permutation Feature Importance

To perform the model inspection technique known as Permutation Feature Importance you will use scikit-learn's built-in function `permutation_importance`.

You will create a function that given a classifier, features and labels computes the feature importance for every feature:

```
In [5]: from sklearn.inspection import permutation_importance

        def feature_importance(clf, X, y, top_limit=None):

            # Retrieve the Bunch object after 50 repeats
            # n_repeats is the number of times that each feature was permuted to compute the final score
            bunch = permutation_importance(clf, X, y,
                                         n_repeats=50, random_state=42)

            # Average feature importance
            imp_means = bunch.importances_mean

            # List that contains the index of each feature in descending order of importance
            ordered_imp_means_args = np.argsort(imp_means)[::-1]

            # If no limit print all features
            if top_limit is None:
                top_limit = len(ordered_imp_means_args)

            # Print relevant information
            for i, _ in zip(ordered_imp_means_args, range(top_limit)):
                name = data.feature_names[i]
                imp_score = imp_means[i]
                imp_std = bunch.importances_std[i]
                print(f"Feature {name} with index {i} has an average importance score of {imp_score:.3f} +/- {imp_std:.3f}")
```

The importance score is computed in a way that higher values represent better predictive power. To know exactly how it is computed check out this [link](#).

Now use the `feature_importance` function on the Random Forest classifier and the train set:

```
In [6]: feature_importance(rf_clf, X_train, y_train)

Feature flavanoids with index 6 has an average importance score of 0.227 +/- 0.025

Feature proline with index 12 has an average importance score of 0.142 +/- 0.019

Feature color_intensity with index 9 has an average importance score of 0.112 +/- 0.023

Feature od280/od315_of_diluted_wines with index 11 has an average importance score of 0.007 +/- 0.005
```

Feature total_phenols with index 5 has an average importance score of 0.003 +/- 0.004
Feature malic_acid with index 1 has an average importance score of 0.002 +/- 0.004
Feature proanthocyanins with index 8 has an average importance score of 0.002 +/- 0.003
Feature hue with index 10 has an average importance score of 0.002 +/- 0.003
Feature nonflavanoid_phenols with index 7 has an average importance score of 0.000 +/- 0.000
Feature magnesium with index 4 has an average importance score of 0.000 +/- 0.000
Feature alkalinity_of_ash with index 3 has an average importance score of 0.000 +/- 0.000
Feature ash with index 2 has an average importance score of 0.000 +/- 0.000
Feature alcohol with index 0 has an average importance score of 0.000 +/- 0.000

Looks like many of the features have a fairly low importance score. This points that the predictive power of this dataset is condensed in a few features.

However it is important to notice that this process was done for the training set, so this feature importance does NOT have into account if the feature might help with the generalization power of the model.

To check this, repeat the process for the test set:

In [7]:

```
feature_importance(rf_clf, X_test, y_test)
```

Feature flavanoids with index 6 has an average importance score of 0.202 +/- 0.047
Feature proline with index 12 has an average importance score of 0.143 +/- 0.042
Feature color_intensity with index 9 has an average importance score of 0.112 +/- 0.043
Feature alcohol with index 0 has an average importance score of 0.024 +/- 0.017
Feature magnesium with index 4 has an average importance score of 0.021 +/- 0.015
Feature od280/od315_of_diluted_wines with index 11 has an average importance score of 0.015 +/- 0.018
Feature hue with index 10 has an average importance score of 0.013 +/- 0.018
Feature total_phenols with index 5 has an average importance score of 0.002 +/- 0.016
Feature nonflavanoid_phenols with index 7 has an average importance score of 0.000 +/- 0.000
Feature alkalinity_of_ash with index 3 has an average importance score of 0.000 +/- 0.000
Feature malic_acid with index 1 has an average importance score of -0.002 +/- 0.017
Feature ash with index 2 has an average importance score of -0.003 +/- 0.008
Feature proanthocyanins with index 8 has an average importance score of -0.021 +/- 0.020

Notice that the top most important features are the same for both sets. However features such as **alcohol**, which was considered not important for the training set is much more important when using the testing set. This hints that this feature will contribute to the generalization power of the model.

If a feature is deemed as important for the train set but not for the testing, this feature will probably cause the model to overfit.

Re-train the model with the most important features

Now you will re-train the Random Forest classifier with only the top 3 most important features.

In this case they are the same for both sets:

In [8]:

```
print("On TRAIN split:\n")
feature_importance(rf_clf, X_train, y_train, top_limit=3)

print("\nOn TEST split:\n")
feature_importance(rf_clf, X_test, y_test, top_limit=3)
```

On TRAIN split:

Feature flavanoids with index 6 has an average importance score of 0.227 +/- 0.025
Feature proline with index 12 has an average importance score of 0.142 +/- 0.019
Feature color_intensity with index 9 has an average importance score of 0.112 +/- 0.023

On TEST split:

Feature flavanoids with index 6 has an average importance score of 0.202 +/- 0.047
Feature proline with index 12 has an average importance score of 0.143 +/- 0.042
Feature color_intensity with index 9 has an average importance score of 0.112 +/- 0.043

```
In [9]: # Preserve only the top 3 features
X_train_top_features = X_train[:, [6, 9, 12]]
X_test_top_features = X_test[:, [6, 9, 12]]

# Re-train with only these features
rf_clf_top = RandomForestClassifier(n_estimators=10, random_state=42).fit(X_train_top_features, y_train)

# Compute mean accuracy achieved
rf_clf_top.score(X_test_top_features, y_test)
```

Out[9]: 0.9333333333333333

Notice that by using only the 3 most important features the model achieved a mean accuracy even higher than the one using all 13 features.

Remember that the **alcohol** feature was deemed not important in the train split but you had the hypotheses that it had important information for the generalization of the model.

Add this feature and see how the model performs:

```
In [10]: # Preserve only the top 3 features
X_train_top_features = X_train[:, [0, 6, 9, 12]]
X_test_top_features = X_test[:, [0, 6, 9, 12]]

# Re-train with only these features
rf_clf_top = RandomForestClassifier(n_estimators=10, random_state=42).fit(X_train_top_features, y_train)

# Compute mean accuracy achieved
rf_clf_top.score(X_test_top_features, y_test)
```

Out[10]: 1.0

Wow! By adding this additional feature you now get a mean accuracy of 100%! Quite remarkable! Looks like this feature did in fact provided some important information that helped the model do a better job at generalizing.

Try out other classifiers

The process of Permutation Feature Importance is also dependant on the classifier you are using. Since different classifiers follow different rules for classification it is natural to assume they will consider different features to be important or unimportant.

To test this, try out other classifiers:

```
In [11]: from sklearn.svm import SVC
from sklearn.linear_model import Lasso, Ridge
from sklearn.tree import DecisionTreeClassifier

# Select 4 new classifiers
clfs = {"Lasso": Lasso(alpha=0.05),
        "Ridge": Ridge(),
        "Decision Tree": DecisionTreeClassifier(),
        "Support Vector": SVC()}

# Compute feature importance on the test set given a classifier
def fit_compute_importance(clf):
    clf.fit(X_train, y_train)
    print(f"Mean accuracy score on the test set: {clf.score(X_test, y_test)*100:.2f}%\n")
    print(f"Top 4 features when using the test set:\n")
    feature_importance(clf, X_test, y_test, top_limit=4)

# Print results
for name, clf in clfs.items():
    print("====="*20)
    print(f"{name} classifier\n")
    fit_compute_importance(clf)
```

=====

Lasso classifier

Mean accuracy score on the test set: 86.80%

Top 4 features when using the test set:

Feature flavanoids with index 6 has an average importance score of 0.323 +/- 0.055

Feature proline with index 12 has an average importance score of 0.203 +/- 0.035

Feature od280/od315_of_diluted_wines with index 11 has an average importance score of 0.146 +/- 0.030

Feature alcalinity_of_ash with index 3 has an average importance score of 0.038 +/- 0.014

=====

➡ Ridge classifier

✏ Mean accuracy score on the test set: 88.71%

📊 Top 4 features when using the test set:

Feature flavanoids with index 6 has an average importance score of 0.445 +/- 0.071

Feature proline with index 12 has an average importance score of 0.210 +/- 0.035

Feature color_intensity with index 9 has an average importance score of 0.119 +/- 0.029

Feature od280/od315_of_diluted_wines with index 11 has an average importance score of 0.111 +/- 0.026

=====

➡ Decision Tree classifier

✏ Mean accuracy score on the test set: 95.56%

📊 Top 4 features when using the test set:

Feature proline with index 12 has an average importance score of 0.300 +/- 0.056

Feature flavanoids with index 6 has an average importance score of 0.297 +/- 0.061

Feature color_intensity with index 9 has an average importance score of 0.131 +/- 0.037

Feature malic_acid with index 1 has an average importance score of 0.008 +/- 0.012

=====

➡ Support Vector classifier

✏ Mean accuracy score on the test set: 97.78%

📊 Top 4 features when using the test set:

Feature proline with index 12 has an average importance score of 0.069 +/- 0.031

Feature flavanoids with index 6 has an average importance score of 0.061 +/- 0.023

Feature alcohol with index 0 has an average importance score of 0.044 +/- 0.023

Feature ash with index 2 has an average importance score of 0.032 +/- 0.018

Looks like **flavanoids** and **proline** are very important across all classifiers. However there is variability from one classifier to the others on what features are considered the most important ones.

Congratulations on finishing this ungraded lab! Now you should have a clearer understanding of what Permutation Feature Importance is, why it is useful and how to implement this technique using scikit-learn.

Keep it up!